

120A3051**Shreya Idate****Batch: E3**

EXPERIMENT NO. 2

Aim: Data preparation using NumPy and Pandas

- a. Obtain a listing of all records that are outliers according to the any field. Print out a listing of the 10 largest values for that field.
- b. Do the following for the any field.
 - i. Standardize the variable.
 - ii. Identify how many outliers there are and identify the most extreme outlier

Theory:

What is Pandas?

Pandas is a Python library used for working with data sets. It has functions for analyzing, cleaning, exploring, and manipulating data. The name "Pandas" has a reference to both "Panel Data", and "Python Data Analysis" and was created by Wes McKinney in 2008.

Why Use Pandas?

Pandas allows us to analyze big data and make conclusions based on statistical theories. Pandas can clean messy data sets, and make them readable and relevant. Relevant data is very important in data science.

What are Outliers?

An outlier of a dataset is defined as a value that is more than 3 standard deviations from the mean. Removing outliers from a pandas DataFrame removes any rows in the DataFrame which contain an outlier. Outlier calculations are performed separately for each column.

What is nlargest?

DataFrame.nlargest(n, columns, keep='first') [source]

Return the first n rows ordered by columns in descending order. Return the first n rows with the largest values in columns, in descending order. The columns that are not specified are returned as well, but not used for ordering.

What is group by?

DataFrame.groupby(by=None, axis=0, level=None, as_index=True, sort=True, group_keys=True, squeeze=False, observed=False, dropna=True) [source]

Group DataFrame using a mapper or by a Series of columns. A groupby operation involves some

combination of splitting the object, applying a function, and combining the results. This can be used to group large amounts of data and compute operations on these groups.

What is mean?

`DataFrame.mean(axis=NoDefault.no_default, skipna=True, level=None, numeric_only=None, **kwargs)` [source]

Return the mean of the values over the requested axis.

Program:

```

import pandas as pd
import numpy as np
import plotly.express as px
[2] ✓ 0.1s Python

df = pd.read_csv(r"C:\Users\exam\Desktop\120A3051\uber.csv")

df.head(10)
[8] ✓ 0.4s Python

...   Unnamed: 0   key   fare_amount   pickup_datetime   pickup_longitude   pickup_latitude   dropoff_longitude   dropoff_latitude   passenger_count
0    24238194  2015-05-07 19:52:06.0000003   7.5  2015-05-07 19:52:06 UTC   -73.999817   40.738354   -73.999512   40.723217   1
1    27835199  2009-07-17 20:04:56.0000002   7.7  2009-07-17 20:04:56 UTC   -73.994355   40.728225   -73.994710   40.750325   1
2    44984355  2009-08-24 21:45:00.00000061  12.9  2009-08-24 21:45:00 UTC   -74.005043   40.740770   -73.962565   40.772647   1
3    25894730  2009-06-26 08:22:21.0000001   5.3  2009-06-26 08:22:21 UTC   -73.976124   40.790844   -73.965316   40.803349   3
4    17610152  2014-08-28 17:47:00.000000188  16.0  2014-08-28 17:47:00 UTC   -73.925023   40.744085   -73.973082   40.761247   5
5    44470845  2011-02-12 02:27:09.0000006   4.9  2011-02-12 02:27:09 UTC   -73.969019   40.755910   -73.969019   40.755910   1
6    48725865  2014-10-12 07:04:00.0000002   24.5  2014-10-12 07:04:00 UTC   -73.961447   40.693965   -73.871195   40.774297   5
7    44195482  2012-12-11 13:52:00.00000029  2.5   2012-12-11 13:52:00 UTC   0.000000   0.000000   0.000000   0.000000   1
8    15822268  2012-02-17 09:32:00.00000043  9.7   2012-02-17 09:32:00 UTC   -73.975187   40.745767   -74.002720   40.743537   1
9    50611056  2012-03-29 19:06:00.000000273  12.5  2012-03-29 19:06:00 UTC   -74.001065   40.741787   -73.963040   40.775012   1

```



```

df = df.drop(columns=['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude'])
df
[9] ✓ 0.8s Python

...   Unnamed: 0   key   fare_amount   pickup_datetime   passenger_count
0    24238194  2015-05-07 19:52:06.0000003   7.5  2015-05-07 19:52:06 UTC   1
1    27835199  2009-07-17 20:04:56.0000002   7.7  2009-07-17 20:04:56 UTC   1
2    44984355  2009-08-24 21:45:00.00000061  12.9  2009-08-24 21:45:00 UTC   1
3    25894730  2009-06-26 08:22:21.0000001   5.3  2009-06-26 08:22:21 UTC   3
4    17610152  2014-08-28 17:47:00.000000188  16.0  2014-08-28 17:47:00 UTC   5
...
199995  42598914  2012-10-28 10:49:00.00000053  3.0  2012-10-28 10:49:00 UTC   1
199996  16382965  2014-03-14 01:09:00.0000008   7.5  2014-03-14 01:09:00 UTC   1
199997  27804658  2009-06-29 00:42:00.00000078  30.9  2009-06-29 00:42:00 UTC   2
199998  20259894  2015-05-20 14:56:25.0000004   14.5  2015-05-20 14:56:25 UTC   1
199999  11951496  2010-05-15 04:08:00.00000076  14.1  2010-05-15 04:08:00 UTC   1

```



200000 rows × 5 columns

```
[11] df.describe[['fare_amount','passenger_count']]  
✓ 0.1s  
...  




|       | fare_amount   | passenger_count |
|-------|---------------|-----------------|
| count | 200000.000000 | 200000.000000   |
| mean  | 11.359955     | 1.684535        |
| std   | 9.901776      | 1.385997        |
| min   | -52.000000    | 0.000000        |
| 25%   | 6.000000      | 1.000000        |
| 50%   | 8.500000      | 1.000000        |
| 75%   | 12.500000     | 2.000000        |
| max   | 499.000000    | 208.000000      |


```


With x=fare_amount



With x=passenger_count



Scatter with x n y



Finding outliers using statistical methods

```

def find_outliers(df):
    q1 = df.quantile(0.25)
    q3 = df.quantile(0.75)
    IQR = q3 - q1
    outliers = df[(df<(q1-1.5*IQR)) | (df>(q3+1.5*IQR))]
    return outliers

```

[21] ✓ 0.2s

```

outliers = find_outliers(df['fare_amount'])

print('No. of outliers: ' + str(len(outliers)))
print('Maximum outlier value: ' + str(outliers.max()))
print('Minimum outlier value: ' + str(outliers.min()))

outliers

```

[22] ✓ 0.8s

... No. of outliers: 17167
 Maximum outlier value: 499.0
 Minimum outlier value: -52.0

6	24.50
30	25.70
34	39.50
39	29.00
48	56.80
	...
199976	49.70
199977	43.50
199982	57.33
199985	24.00
199997	30.90

Name: fare_amount, Length: 17167, dtype: float64

```
outliers.sort_values(ascending=False).head(10)
```

[26] ✓ 0.3s

```
... 170081    499.00
    4292     350.00
    185325     275.00
    71715      250.00
    197493     230.00
    29261      220.00
    23682      206.38
    196647      200.00
    184901      196.00
    33911      192.33
Name: fare_amount, dtype: float64
```

```
passenger_outliers = find_outliers(df['passenger_count'])
```

```
print('No. of outliers: ' + str(len(outliers)))
print('Maximum outlier value: ' + str(outliers.max()))
print('Minimum outlier value: ' + str(outliers.min()))
```

```
passenger_outliers
```

[28] ✓ 0.1s

```
... No. of outliers: 22557
Maximum outlier value: 208
Minimum outlier value: 4
```

```
4      5
6      5
12     5
24     5
29     5
...
199958  5
199959  5
199962  4
199969  5
199985  5
```

```
Name: passenger_count, Length: 22557, dtype: int64
```

```
[29]    passenger_outliers.sort_values(ascending=False).head(10)
✓ 0.2s
...
... 113038    208
123723      6
164226      6
164250      6
123654      6
78604       6
78704       6
78713       6
36785       6
164344      6
Name: passenger_count, dtype: int64
```

Working with outliers using statistical methods

Drop unnecessary columns

```
[31] df=df.drop(columns=[ 'Unnamed: 0','key','pickup_datetime']))
✓ 0.2s
▷ ▾ df
[32] ✓ 0.3s
...
...      fare_amount  passenger_count
0           7.5              1
1           7.7              1
2          12.9              1
3           5.3              3
4          16.0              5
...
199995     3.0              1
199996     7.5              1
199997    30.9              2
199998    14.5              1
199999    14.1              1
200000 rows × 2 columns
```

STANDARDIZE THE DATASET

```
pip install -U scikit-learn scipy matplotlib
```

```
[34]     from sklearn.preprocessing import StandardScaler
      ✓  3.8s

[35]     scale= StandardScaler()
      scaled_data = scale.fit_transform(df)
      print(scaled_data)
      ✓  0.3s
...
[[[-0.3898255 -0.49389496]
 [-0.36962706 -0.49389496]
 [ 0.15553256 -0.49389496]
 ...
 [ 1.97339277  0.22760936]
 [ 0.31712013 -0.49389496]
 [ 0.27672324 -0.49389496]]]
```

```
[36]     type(scaled_data)
      ✓  0.3s
...
numpy.ndarray
```

```
[38]     new_df = pd.DataFrame(scaled_data, columns = ['fare_amount', 'passenger_count'])
      print(new_df)
      print(type(new_df))
      ✓  0.3s
...
       fare_amount  passenger_count
0           -0.389826      -0.493895
1           -0.369627      -0.493895
2            0.155533      -0.493895
3           -0.612008      0.949114
4            0.468608      2.392122
...
       ...
199995      -0.844291      -0.493895
199996      -0.389826      -0.493895
199997      1.973393      0.227609
199998      0.317120      -0.493895
199999      0.276723      -0.493895

[200000 rows x 2 columns]
<class 'pandas.core.frame.DataFrame'>
```

Conclusion: Successful Data preparation using NumPy and Pandas.